

Zhennan Shen

1641225799szn@gmail.com · zshenbl@connect.ust.hk · zhennan0521.github.io · Google Scholar · GitHub · LinkedIn

RESEARCH SUMMARY

I work on post-training for long-horizon coding agents, focusing on scalable data synthesis and reinforcement learning. My experience covers agent data pipelines, supervised fine-tuning (SFT), RL experiments, and agent evaluation, and I have contributed to the Kimi K2.5, JoyAI-LLM Flash, and MiniMax M3.x model series.

EDUCATION

The Hong Kong University of Science and Technology Sep 2026 – Present
Ph.D. in Computer Science and Engineering · Hong Kong
Advisor: Prof. Yi R. (May) Fung · Tang Xiao'ou Scholarship

Shanghai Jiao Tong University Sep 2021 – Jun 2025
B.Eng. in Computer Science and Engineering · Shanghai
Major GPA: 3.9/4.3 · Zhiyuan Honors Scholarship · Outstanding Graduate

INDUSTRY EXPERIENCE

MiniMax Jun 2026 – Present
Research Intern · Shanghai

- Post-training for long-horizon coding agents on the upcoming M3.1 model, covering SWE-Marathon, NL2Repo, and Frontier-SWE benchmarks.
- Built data synthesis pipelines, ran agent evaluation, validated SFT data quality, and ran RL training experiments.

JD.com · JoyAI Dec 2025 – Apr 2026
Research Intern · Beijing

- Agent data synthesis for mid-training and agent post-training of JoyAI-LLM Flash, covering coding and general agents.
- Contributed to RL infrastructure and training experiments.

Moonshot AI · Kimi, Agent Team Mar 2025 – Dec 2025
Research Intern · Beijing

- Post-training for the computer-use agent (CUA) of Kimi K2.5, covering scalable data synthesis and training iterations, with a focus on OSWorld.

SELECTED PUBLICATIONS & TECHNICAL REPORTS

† indicates first author or co-first author.

Kimi K2.5: Visual Agentic Intelligence. Kimi Team. Technical Report, 2025.

JoyAI-LLM Flash: Advancing Mid-Scale LLMs with Token Efficiency. Technical Report, 2026.

† **AgentTrek: Agent Trajectory Synthesis via Guiding Replay with Web Tutorials.** ICLR 2025 (Spotlight).

† **Evaluating Parameter Efficient Methods for RLVR.** ICML 2026.

† **On the Geometry of On-Policy Distillation.** Preprint, 2026.

OsWorld-Verified. NeurIPS 2024 (Datasets & Benchmarks).

OpenCUA: Open Foundations for Computer-Use Agents. NeurIPS 2025 (Spotlight).

CATP: Contextually Adaptive Token Pruning for Efficient and Enhanced Multimodal In-Context Learning. AAAI 2026 (Oral).

SciEval: A Multi-Level Large Language Model Evaluation Benchmark for Scientific Research. AAAI 2024.

RESEARCH EXPERIENCE

X-LANCE Lab, SJTU | Evaluating LLM capabilities on GUI-agent tasks and building high-quality evaluation resources.

HKU NLP Group | Data synthesis and training pipelines for GUI agents in interactive environments.

Moonshot AI, Agent Team | Built synthetic data and environment pipelines for GUI agents, and curated SFT data for Kimi K2.5 post-training iterations.

JD JoyAI | Built synthetic data pipelines and RL framework workflows for tool-use and coding agents, and validated pretraining data with ablation studies.

MiniMax | Post-training for long-horizon coding agents in the upcoming M3.1 model, focusing on SWE-Marathon, NL2Repo, and Frontier-SWE.

HONORS & AWARDS

- Tang Xiao'ou Scholarship, The Hong Kong University of Science and Technology.
- Zhiyuan Honors Scholarship; Outstanding Graduate, Shanghai Jiao Tong University.